

assignment1

April 29, 2023

```
[46]: # CS614 Assignment #1
      ## Vision
      ### Author: Cory W. Mauer
```

0.1 Pitch:

My Original idea for this project was inspired when scrolling through instagram looking at popular tattoo artists work. The challenge with this is I am often looking for ideas for work, without wanting to make a direct copy of something that already exists. Based on this idea I theorized the idea of building web site that enabled users to view autogenerated art images as a means for inspiration toward their own artwork.

0.2 Data source:

During my initial investigation of data sets to use for this project I initial was looking to either track down an existing data set of tattoo related art, however unsurprisingly I was not able to find any.

Following that discovery I decided I wanted to track down more general art related data. While I was able to find some, most had inconsistent large and/or inconsistent image sizing which made the process more difficult.

For this project I finally decided to leverage the [Ciphar-10 data set](https://www.tensorflow.org/datasets/catalog/cifar10) (<https://www.tensorflow.org/datasets/catalog/cifar10>) as it provided my with an easily accessible large, color, and tagged dataset leveraging a wide variety of classes. Originally I expected very poor performance when using this data set, however after using it I was happily surprised by the results. They say that art imitates life, and in this case the combination of data from seemingly randomly classes tends to actually produce some very interesting images.

here are same examples from the data set with class labels


```
[47]: import os
      os.environ['TF_CPP_MIN_LOG_LEVEL'] = '3'
      import tensorflow as tf
      from tensorflow.keras.datasets import cifar10
      from matplotlib import pyplot as plt

      (train_images, train_labels), (_, _) = cifar10.load_data()
```

```

figure_size = {"length": 5, "width":5}
fig = plt.figure(figsize=(figure_size["length"],figure_size["width"]))
for i in range(figure_size["length"]* figure_size["width"]):
    plt.subplot(figure_size["length"], figure_size["width"], i+1)
    plt.title(f"class: {train_labels[i]}", fontsize=6)
    plt.imshow((train_images[i, :, :, :]))
    plt.axis('off')

```



Given more time and compute resources, I would like to build or aquire larger and more targeted datasets. With these larger data sets I would hope to build higher resolution artwork that can better target specific use cases.

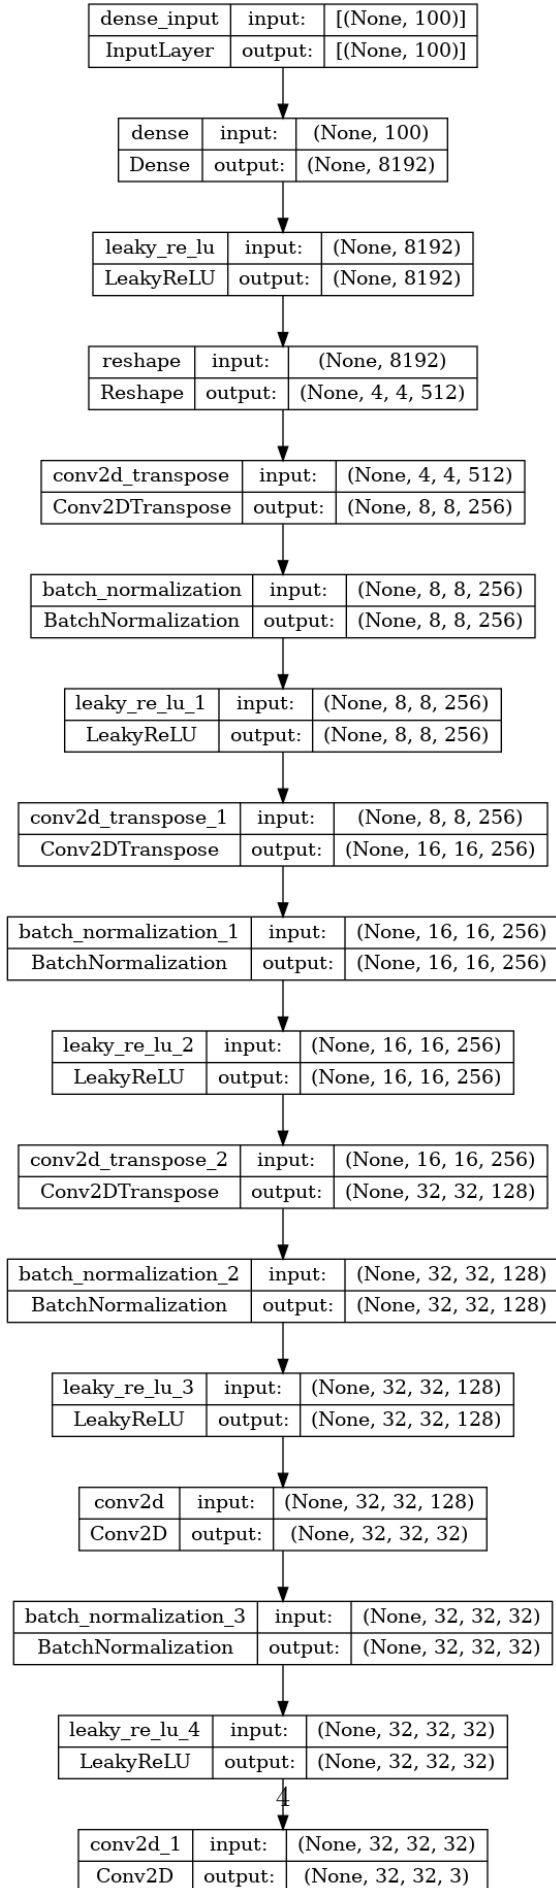
0.3 Model and Data Justification

For this task I chose to use a GAN for the image generation. Prior to this project I had read about GAN's at a high level but never actually implemented one. Given more time and compute resources I would have liked to have explored a more cutting edge model such as a diffusion model.

During my GAN development, I originally started with a simple GAN consisting of convolutional transpose layers with relatively few chanel.

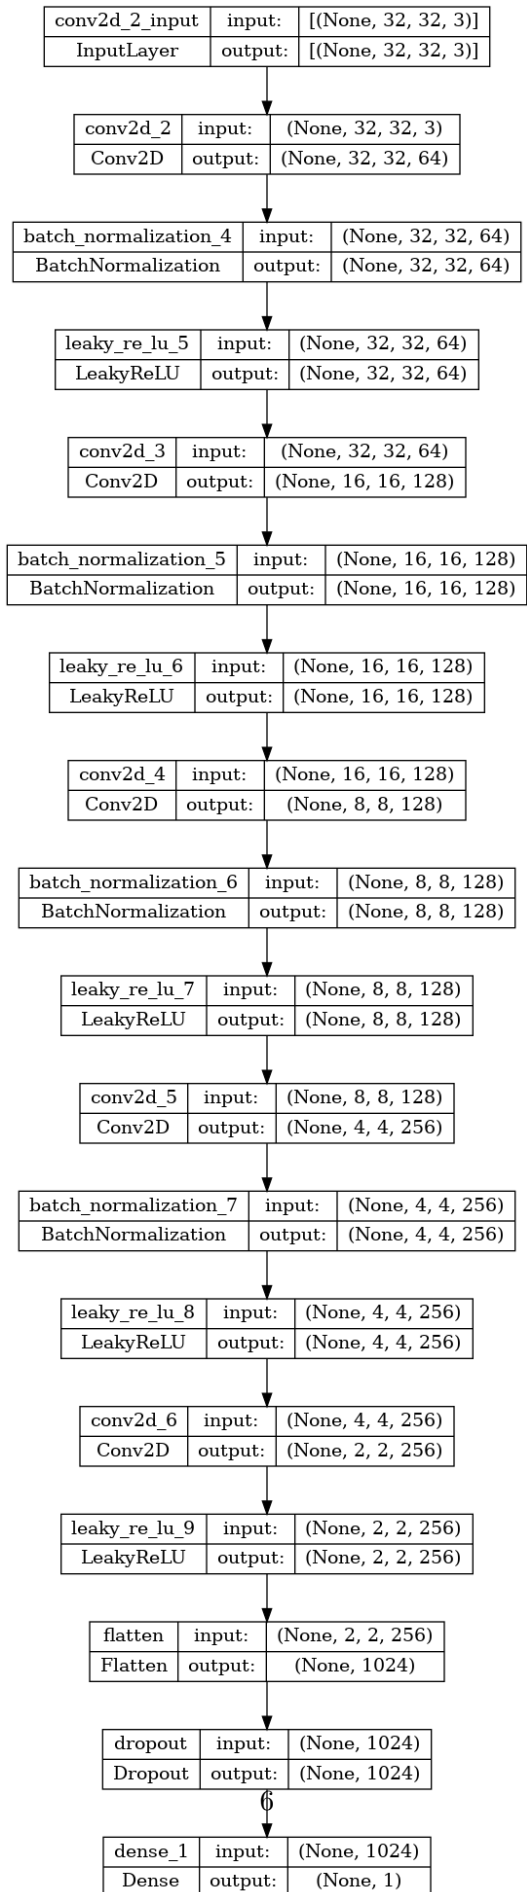
best gan generator

```
[48]: from IPython.display import display, Image
      display(Image(filename="./model_plots/generator_plot.png", height=400,
      ↪width=400))
```



best gan discriminator

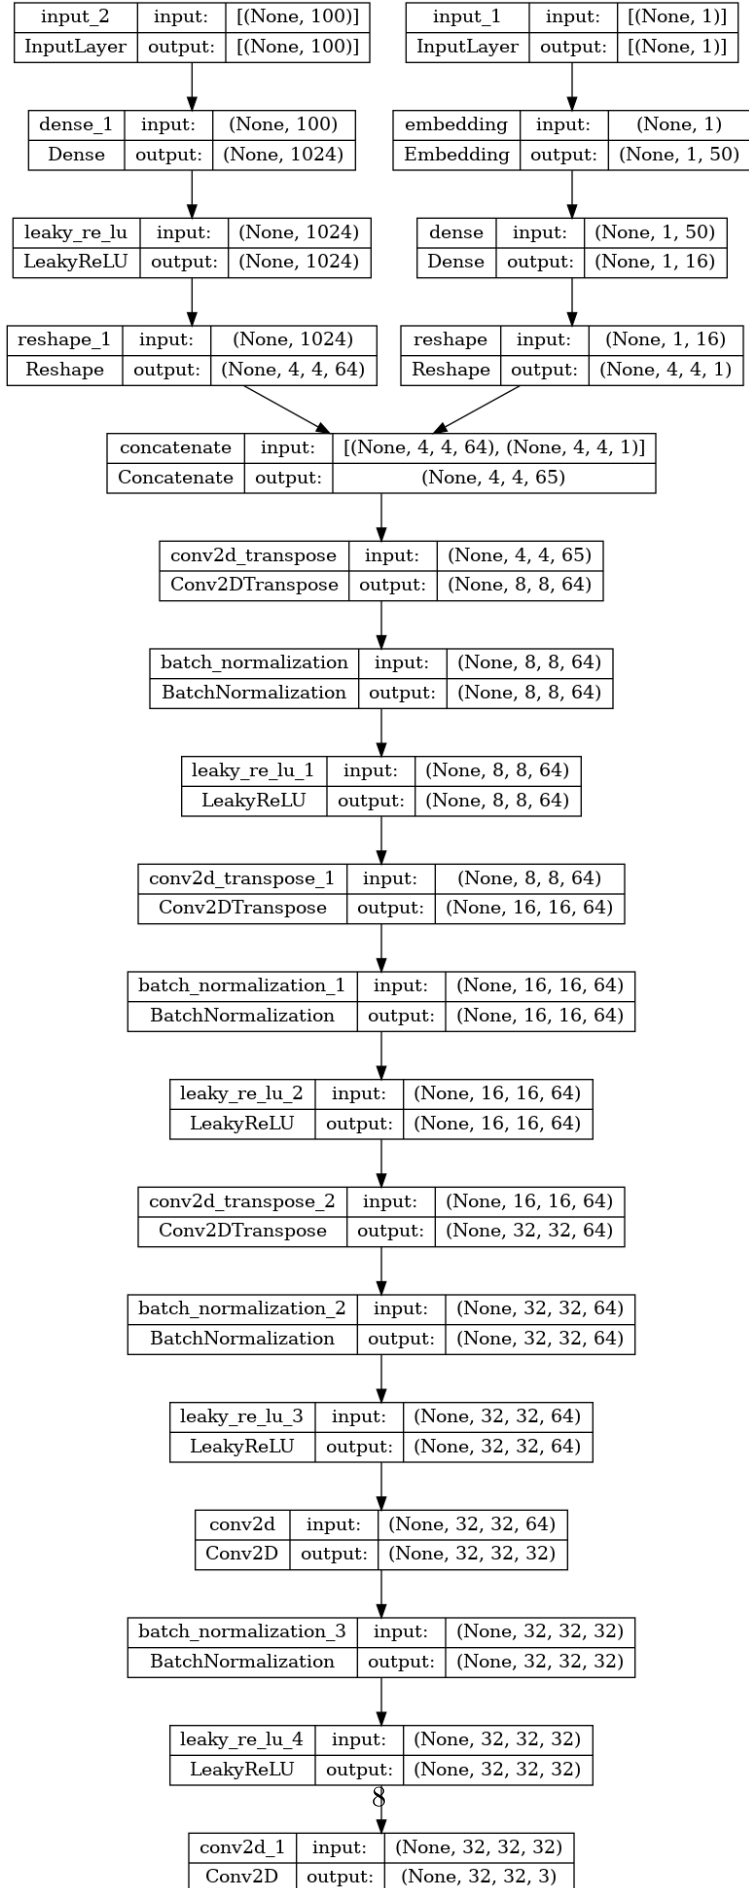
```
[49]: display(Image(filename="./model_plots/discriminator_plot.png", height=400, ↵  
      ↪width=400))
```



I did start exploring using classification in effort to create a conditional GAN, but didn't have great success with training.

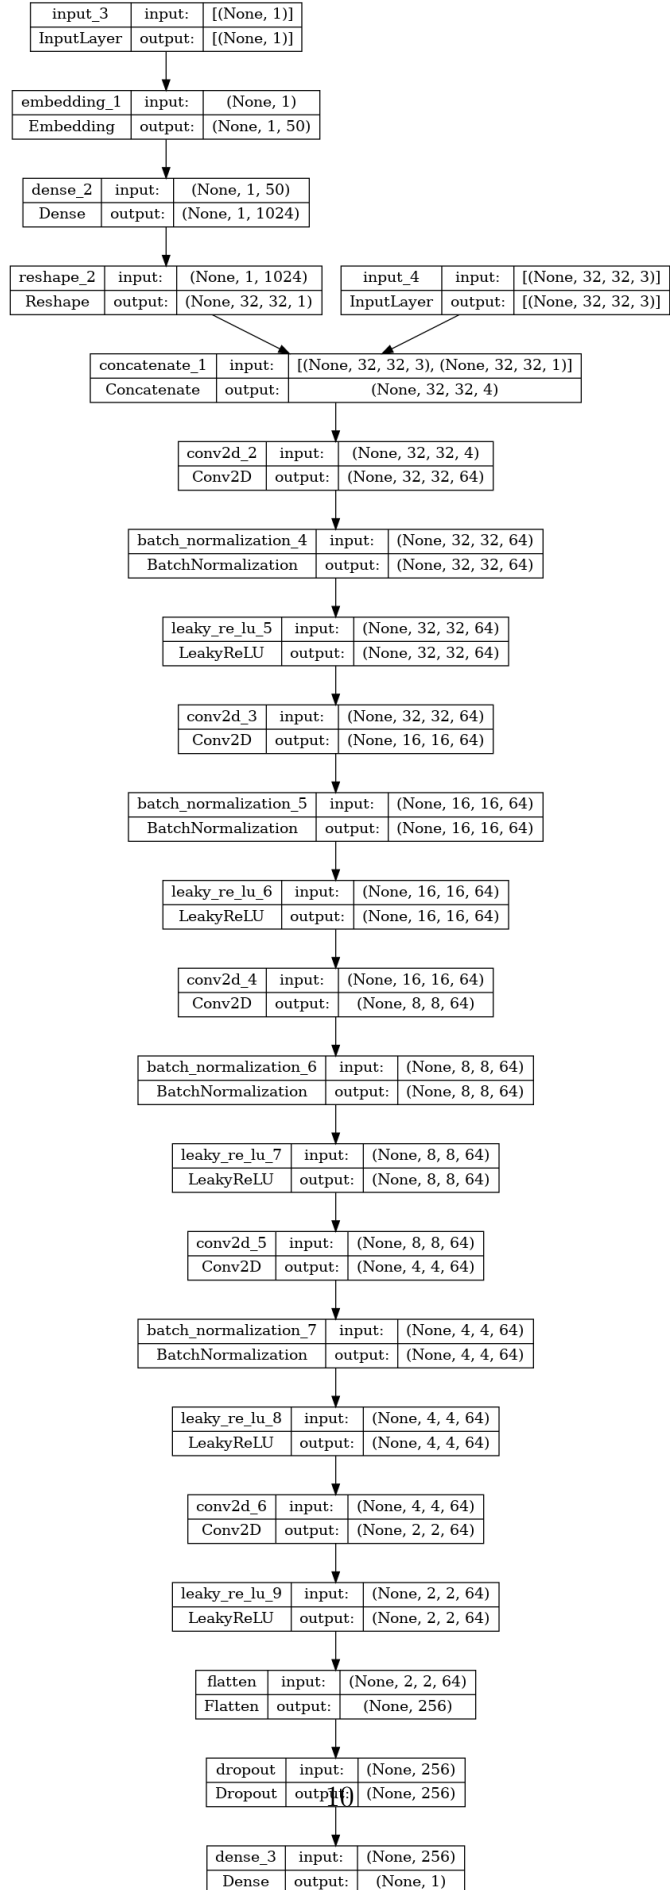
cgan generator

```
[50]: display(Image(filename="./model_plots/cgan_gen_plot.png", height=400,  
    ↪width=400))
```



cgan discriminator

```
[51]: display(Image(filename="./model_plots/cgan_discriminator_plot.png", height=400, width=400))
```



0.4 Commented examples

The following provides some commented code snippets of the model. for more details see the training script `cifar10_gan.py`

```
[52]: from tensorflow.keras import layers
      # Generator model
      def make_generator_model():
          model = tf.keras.Sequential()
          n_nodes = [4, 4, 512]
          model.add(layers.Dense(n_nodes[0] * n_nodes[1] * n_nodes[2], input_dim=100))
          model.add(layers.LeakyReLU(alpha=0.2))
          model.add(layers.Reshape(n_nodes))
          model.add(layers.Conv2DTranspose(256, (4,4), strides=(2,2), padding='same'))
          model.add(layers.BatchNormalization())
          model.add(layers.LeakyReLU(alpha=0.2))
          model.add(layers.Conv2DTranspose(256, (4,4), strides=(2,2), padding='same'))
          model.add(layers.BatchNormalization())
          model.add(layers.LeakyReLU(alpha=0.2))
          model.add(layers.Conv2DTranspose(128, (4,4), strides=(2,2), padding='same'))
          model.add(layers.BatchNormalization())
          model.add(layers.LeakyReLU(alpha=0.2))
          model.add(layers.Conv2D(32, (3,3), padding='same'))
          model.add(layers.BatchNormalization())
          model.add(layers.LeakyReLU(alpha=0.2))
          model.add(layers.Conv2D(3, (3,3), activation='tanh', padding='same'))
          return model

      generator = make_generator_model()
      print(f"generator \n {generator.summary()}")
```

Model: "sequential_8"

Layer (type)	Output Shape	Param #
dense_8 (Dense)	(None, 8192)	827392
leaky_re_lu_40 (LeakyReLU)	(None, 8192)	0
reshape_4 (Reshape)	(None, 4, 4, 512)	0
conv2d_transpose_12 (Conv2D Transpose)	(None, 8, 8, 256)	2097408
batch_normalization_32 (Batch Normalization)	(None, 8, 8, 256)	1024

```

chNormalization)

leaky_re_lu_41 (LeakyReLU) (None, 8, 8, 256) 0

conv2d_transpose_13 (Conv2D Transpose) (None, 16, 16, 256) 1048832

batch_normalization_33 (Batch Normalization) (None, 16, 16, 256) 1024

leaky_re_lu_42 (LeakyReLU) (None, 16, 16, 256) 0

conv2d_transpose_14 (Conv2D Transpose) (None, 32, 32, 128) 524416

batch_normalization_34 (Batch Normalization) (None, 32, 32, 128) 512

leaky_re_lu_43 (LeakyReLU) (None, 32, 32, 128) 0

conv2d_28 (Conv2D) (None, 32, 32, 32) 36896

batch_normalization_35 (Batch Normalization) (None, 32, 32, 32) 128

leaky_re_lu_44 (LeakyReLU) (None, 32, 32, 32) 0

conv2d_29 (Conv2D) (None, 32, 32, 3) 867

=====
Total params: 4,538,499
Trainable params: 4,537,155
Non-trainable params: 1,344
-----
generator
None

```

```

[53]: # Discriminator model
def make_discriminator_model():
    model = tf.keras.Sequential()
    model.add(layers.Conv2D(64, (3,3), padding='same', input_shape=(32,32,3)))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.Conv2D(128, (3,3), strides=(2,2), padding='same'))
    model.add(layers.BatchNormalization())
    model.add(layers.LeakyReLU(alpha=0.2))
    model.add(layers.Conv2D(128, (3,3), strides=(2,2), padding='same'))

```

```

model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU(alpha=0.2))
model.add(layers.Conv2D(256, (3,3), strides=(2,2), padding='same'))
model.add(layers.BatchNormalization())
model.add(layers.LeakyReLU(alpha=0.2))
model.add(layers.Conv2D(256, (3,3), strides=(2,2), padding='same'))
model.add(layers.LeakyReLU(alpha=0.2))
model.add(layers.Flatten())
model.add(layers.Dropout(0.25))
model.add(layers.Dense(1, activation='sigmoid'))
return model

discriminator = make_discriminator_model()
print(f"discriminator \n {discriminator.summary()}")

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
conv2d_30 (Conv2D)	(None, 32, 32, 64)	1792
batch_normalization_36 (Batch Normalization)	(None, 32, 32, 64)	256
leaky_re_lu_45 (LeakyReLU)	(None, 32, 32, 64)	0
conv2d_31 (Conv2D)	(None, 16, 16, 128)	73856
batch_normalization_37 (Batch Normalization)	(None, 16, 16, 128)	512
leaky_re_lu_46 (LeakyReLU)	(None, 16, 16, 128)	0
conv2d_32 (Conv2D)	(None, 8, 8, 128)	147584
batch_normalization_38 (Batch Normalization)	(None, 8, 8, 128)	512
leaky_re_lu_47 (LeakyReLU)	(None, 8, 8, 128)	0
conv2d_33 (Conv2D)	(None, 4, 4, 256)	295168
batch_normalization_39 (Batch Normalization)	(None, 4, 4, 256)	1024
leaky_re_lu_48 (LeakyReLU)	(None, 4, 4, 256)	0

conv2d_34 (Conv2D)	(None, 2, 2, 256)	590080
leaky_re_lu_49 (LeakyReLU)	(None, 2, 2, 256)	0
flatten_4 (Flatten)	(None, 1024)	0
dropout_4 (Dropout)	(None, 1024)	0
dense_9 (Dense)	(None, 1)	1025

```
=====
Total params: 1,111,809
Trainable params: 1,110,657
Non-trainable params: 1,152
-----
discriminator
None
```

0.5 Testing


Initial testing using these gan models focused on training data from a single class from the cifar data set. In this case we will demonstrate the model being trained with automobile images. Notice the quality of the relatively high quality of the images that make it easy to tell that a car. The homogenous nature of the images all being cars allow for a model that clearly produces car images.

```
[54]: generator = tf.keras.models.load_model("./car_generator")
figure_size = {"length": 5, "width": 5}
number_of_generated_images = figure_size['length'] * figure_size['width']
noise = tf.random.normal([number_of_generated_images, 100])
generated_images = generator.predict(noise)
generated_images = (generated_images * 127.5 + 127.5).astype(int)

fig = plt.figure(figsize=(figure_size["length"], figure_size["width"]), dpi= 100)
for i in range(figure_size["length"]* figure_size["width"]):
    plt.subplot(figure_size["length"], figure_size["width"], i+1)
    plt.imshow((generated_images[i, :, :, :]))
    plt.axis('off')
```

WARNING:tensorflow:No training configuration found in save file, so the model was *not* compiled. Compile it manually.

1/1 [=====] - 0s 134ms/step




After this I decided look into how I could make artistic looking images. I found that by training my models with all classes of image from the cifar dataset I was able to create images that mix features of each of the classes and make unique but familiar images.

```
[55]: generator = tf.keras.models.load_model("./gan10_generator")
figure_size = {"length": 10, "width":10}
number_of_generated_images = figure_size['length'] * figure_size['width']
noise = tf.random.normal([number_of_generated_images, 100])
generated_images = generator.predict(noise)
generated_images = (generated_images * 127.5 + 127.5).astype(int)

fig = plt.figure(figsize=(figure_size["length"],figure_size["width"]), dpi= 800)
for i in range(figure_size["length"]* figure_size["width"]):
    plt.subplot(figure_size["length"], figure_size["width"], i+1)
    plt.imshow((generated_images[i, :, :, :]))
    plt.axis('off')
```

WARNING:tensorflow:No training configuration found in save file, so the model was *not* compiled. Compile it manually.

4/4 [=====] - 0s 20ms/step



0.6 Code and instructions to run it

Code can be found on my github at [cs614](https://github.com/cwma86/cs614)(<https://github.com/cwma86/cs614>) The model training code exists in `./cifar10_gan.py`. the pre-trained model can be found at `gan10_generator` and `gan10_discriminator`

to add more training epoch run

```
./cifar10_gan.py --gen gan10_generator --disc gan10_discriminator
```

to see super resolved version of the generated images

```
./super_resolution.py gan10_generator
```